

# Data Exploring and Analysis

Nowadays, massive data is collected daily and distributed over various channels. This requires efficient and flexible data analysis tools. Python's open source Pandas library fills that gap and deals with three different data structures: series, data frames, and panels. A *series* is a one-dimensional data structure such as a dictionary, array, list, tuple, and so on. A *data frame* is a two-dimensional data structure with heterogeneous data types, i.e., tabular data. A *panel* refers to a three-dimensional data structure such as a three-dimensional array. It should be clear that the higher-dimensional data structure is a container of its lower-dimensional data structure. In other words, a panel is a container of a data frame, and a data frame is a container of a series.

## Series Data Structures

As mentioned earlier, a series is a sequence of one-dimensional data such as a dictionary, list, array, tuple, and so on.

## Creating a Series

Pandas provides a `Series()` method that is used to create a series structure. A series structure of size  $n$  should have an index of length  $n$ . By default Pandas creates indices starting at 0 and ending with  $n-1$ . A Pandas series can be created using the constructor `pandas.Series (data, index, dtype, copy)` where `data` could be an array, constant, list, etc. The series index should be unique and hashable with length  $n$ , while `dtype` is a data type that could be explicitly declared or inferred from the received data. Listing 6-1 creates a series with a default index and with a set index.

### **Listing 6-1.** Creating a Series

```
In [5]: import pandas as pd
import numpy as np
data = np.array(['O','S','S','A'])
S1 = pd.Series(data) # without adding index
S2 = pd.Series(data,index=[100,101,102,103]) # with
adding index print (S1) print ("\n") print (S2)
0    O
1    S
2    S
3    A
dtype: object

100  O
101  S
102  S
103  A
dtype: object
```

```
In [40]:import pandas as pd
import numpy as np
my_series2 = np.random.randn(5, 10)
print ("\nmy_series2\n", my_series2)
```

This is the output of creating a series of random values of 5 rows and 10 columns.

```
my_series2
[[ 0.08590877  0.59702919 -1.29330859 -1.42021041 -0.09535271  0.09058623
 -1.14191133 -0.84699991  0.94028641  1.79400706]
 [ 0.50645411 -0.37674882 -1.16751734 -1.24061761  0.03981985  0.13478382
  0.76132521 -0.40671662 -0.7484758  0.30420489]
 [-0.66951224 -1.19373055  1.86446782  1.43047631 -0.06302096  0.49239499
 -0.48208329 -1.9805521  -0.73735706 -1.03152802]
 [-0.79181088  1.02769491 -1.27216885  0.20320462  0.19385809 -0.51614599
 -0.66898612 -0.60962025 -1.43724096 -0.22663712]
 [ 1.14193093 -0.8842498  0.22409272 -0.29599594  1.1917404  1.09016684
  1.87701454  1.08452103 -1.49587483 -0.31887386]]
```

As mentioned earlier, you can create a series from a dictionary; Listing 6-2 demonstrates how to create an index for a data series.

### **Listing 6-2.** Creating an Indexed Series

```
In [6]: import pandas as pd
import numpy as np
data = {'X' : 0., 'Y' : 1., 'Z' : 2.}
SERIES1 = pd.Series(data)
print (SERIES1)
X 0.0
Y 1.0
Z 2.0
dtype: float64

In [7]: import pandas as pd
import numpy as np
data = {'X' : 0., 'Y' : 1., 'Z' : 2.}
SERIES1 = pd.Series(data,index=['Y','Z','W','X'])
print (SERIES1)
Y 1.0
```

```
Z 2.0
W NaN
X 0.0
dtype: float64
```

If you can create series data from a scalar value as shown in Listing 6-3, then an index is mandatory, and the scalar value will be repeated to match the length of the given index.

**Listing 6-3.** Creating a Series Using a Scalar

```
In [9]: # Use sclara to create a series
import pandas as pd
import numpy as np
Series1 = pd.Series(7, index=[0, 1, 2, 3, 4])
print (Series1)
0      7
1      7
2      7
3      7
4      7
dtype: int64
```

## Accessing Data from a Series with a Position

Like lists, you can access a series data via its index value. The examples in Listing 6-4 demonstrate different methods of accessing a series of data. The first example demonstrates retrieving a specific element with index 0. The second example retrieves indices 0, 1, and 2. The third example retrieves the last three elements since the starting index is -3 and moves backward to -2, -1. The fourth and fifth examples retrieve data using the series index labels.

**Listing 6-4.** Accessing a Data Series

```
In [18]: import pandas as pd
        Series1 = pd.Series([1,2,3,4,5],index =
                           ['a','b','c','d','e'])
        print ("Example 1:Retrieve the first element")
        print (Series1[0] )
        print ("\nExample 2:Retrieve the first three element")
        print (Series1[:3])
        print ("\nExample 3:Retrieve the last three element")
        print(Series1[-3:])
        print ("\nExample 4:Retrieve a single element")
        print (Series1['a'])
        print ("\nExample 5:Retrieve multiple elements")
        print (Series1[['a','c','d']])
```

```
Example 1:Retrieve the first element
1
```

```
Example 2:Retrieve the first three element
a    1
b    2
c    3
dtype: int64
```

```
Example 3:Retrieve the last three element
c    3
d    4
e    5
dtype: int64
```

```
Example 4:Retrieve a single element
1
```

```
Example 5:Retrieve multiple elements
a    1
c    3
d    4
dtype: int64
```

## Exploring and Analyzing a Series

Numerous statistical methods can be applied directly on a data series.

Listing 6-5 demonstrates the calculation of mean, max, min, and standard deviation of a data series. Also, the `.describe()` method can be used to give a data description, including quantiles.

### **Listing 6-5.** Analyzing Series Data

```
In [10]: import pandas as pd
import numpy as np
my_series1 = pd.Series([5, 6, 7, 8, 9, 10])
print ("my_series1\n", my_series1)
print ("\n Series Analysis\n ")
print ("Series mean value : ", my_series1.mean()) #
find mean value in a series
print ("Series max value : ",my_series1.max()) #
find max value in a series
print ("Series min value : ",my_series1.min()) #
find min value in a series
print ("Series standard deviation value : ",
my_series1.std()) # find standard deviation
my_series1
0    5
1    6
2    7
3    8
4    9
5   10
dtype: int64
```

## Series Analysis

Series mean value : 7.5

Series max value : 10

Series min value : 5

Series standard deviation value : 1.8708286933869707

In [11]: my\_series1.describe()

```
Out[11]: count      6.000000
         mean       7.500000
         std        1.870829
         min        5.000000
         25%        6.250000
         50%        7.500000
         75%        8.750000
         max        10.000000
         dtype: float64
```

If you copied by reference one series to another, then any changes to the series will adapt to the other one. After copying `my_series1` to `my_series_11`, once you change the indices of `my_series_11`, it reflects back to `my_series1`, as shown in Listing 6-6.

**Listing 6-6.** Copying a Series to Another with a Reference

```
In [17]: my_series_11 = my_series1
         print (my_series1)
         my_series_11.index = ['A', 'B', 'C', 'D', 'E', 'F']
         print (my_series_11)
         print (my_series1)
0        5
1        6
2        7
3        8
```

```

4      9
5     10
dtype: int64
A      5
B      6
C      7
D      8
E      9
F     10
dtype: int64
A      5
B      6
C      7
D      8
E      9
F     10
dtype: int64

```

You can use the `.copy()` method to copy the data set without having a reference to the original series. See Listing 6-7.

**Listing 6-7.** Copying Series Values to Another

```

In [21]: my_series_11 = my_series1.copy()
         print (my_series1)
         my_series_11.index = ['A', 'B', 'C', 'D', 'E', 'F']
         print (my_series_11)
         print (my_series1)
0      5
1      6
2      7
3      8

```



```

4      9
5     10
dtype: int64
A      5
B      6
C      7
D      8
E      9
F     10
dtype: int64
0      5
1      6
2      7
3      8
4      9
5     10
dtype: int64

```

## Operations on a Series

Numerous operations can be implemented on series data. You can check whether an index value is available in a series or not. Also, you can check all series elements against a specific condition, such as if the series value is less than 8 or not. In addition, you can perform math operations on series data directly or via a defined function, as shown in Listing 6-8.

### ***Listing 6-8.*** Operations on Series

```

In [23]: 'F' in my_series_11
Out[23]: True

In [27]: temp = my_series_11 < 8
          temp

```

```

Out[27]: A    True
        B    True
        C    True
        D   False
        E   False
        F   False
        dtype: bool
In [35]: len(my_series_11)

```

```

Out[35]: 6

```

```

In [28]: temp = my_series_11[my_series_11 < 8 ] * 2
        temp

```

```

Out[28]: A    10
        B    12
        C    14
        dtype: int64

```

Define a function to add two series and call the function, like this:

```

In [37]: def AddSeries(x,y):
        for i in range (len(x)):
            print (x[i] + y[i])


In [39]: print ("Add two series\n")
        AddSeries (my_series_11, my_series1)
        Add two series
        10
        12
        14
        16
        18
        20

```

You can visualize data series using the different plotting systems that are covered in Chapter 7. However, Figure 6-1 demonstrates how to get an at-a-glance idea of your series data and graphically explore it via visual plotting diagrams. See Listing 6-9.

**Listing 6-9.** Visualizing Data Series

```
In [49]: import matplotlib.pyplot as plt
        plt.plot(my_series2)
        plt.ylabel('index')
        plt.show()
```




**Figure 6-1.** Line visualization

```
In [54]: from numpy import *
        import math
        import matplotlib.pyplot as plt
        t = linspace(0, 2*math.pi, 400)
```

```
a = sin(t)
b = cos(t)
c = a + b
```

```
In [50]: plt.plot(t, a, 'r') # plotting t, a separately
plt.plot(t, b, 'b') # plotting t, b separately
plt.plot(t, c, 'g') # plotting t, c separately
plt.show()
```

We can add multiple plots to the same canvas as shown in Figure 6-2.



**Figure 6-2.** *Multiplots on the same canvas*

## Data Frame Data Structures

As mentioned earlier, a data frame is a two-dimensional data structure with heterogeneous data types, i.e., tabular data.

## Creating a Data Frame

Pandas can create a data frame using the constructor `pandas.DataFrame`.

`DataFrame(data, index, columns, dtype, copy)`. A data frame can be created from lists, series, dictionaries, Numpy arrays, or other data frames. A Pandas data frame not only helps to store tabular data but also performs arithmetic operations on rows and columns of the data frame. Listing 6-10 creates a data frame from a single list and a list of lists.

### **Listing 6-10.** Creating a Data Frame from a List

```
In [19]: import pandas as pd
        data = [10,20,30,40,50]
        DF1 = pd.DataFrame(data)
        print (DF1)
```

0	10
1	20
2	30
3	40
4	50

```
In [22]: import pandas as pd
        data = [['Ossama',25],['Ali',43],['Ziad',32]]
        DF1 = pd.DataFrame(data,columns=['Name','Age'])
        print (DF1)
```

	Name	Age
0	Ossama	25
1	Ali	43
2	Ziad	32

```
In [21]: import pandas as pd
        data = [['Ossama',25],['Ali',43],['Ziad',32]]
        DF1 = pd.DataFrame(data,columns=['Name','Age'],
                           dtype=float) print (DF1)
```

	Name	Age
0	Ossama	25.0
1	Ali	43.0
2	Ziad	32.0

You can create a data frame from dictionaries or arrays, as shown in Listing 6-11. Also, you can set the data frame indices. However, if you don't set the indices, then the data frame starts with 0 and goes up to  $n-1$ , where  $n$  is the length of the list. Column names are taken by default from the dictionary keys. However, it's possible to set labels for columns as well. The first data frame's `df1` columns are labeled with the dictionary key names; that's why you don't see NaN cases except for the missing value of the project in dictionary 1. While in the second data frame, named `df2`, you change the column name from `Test1` to `Test_1`, and you get NaNs for all the records. This is because of the absence of `Test_1` in the dictionary key of data.

**Listing 6-11.** Creating a DataFrame from a Dictionary

```
In [13]: import pandas as pd
        data = [{'Test1': 10, 'Test2': 20}, {'Test1': 30,
        'Test2': 20, 'Project': 20}]
        # With three column indices, values same as dictionary
        keys
        df1 = pd.DataFrame(data, index=['First', 'Second'],
        columns=['Test2', 'Project' , 'Test1'])

        #With two column indices with one index with another
        name
        df2 = pd.DataFrame(data, index=['First', 'Second'],
        columns=['Project', 'Test_1', 'Test2 '])
        print (df1)
        print ("\n")
        print (df2)
```

	Test2	Project	Test1
First	20	NaN	10
Second	20	20.0	30

	Project	Test_1	Test2
First	NaN	NaN	20
Second	20.0	NaN	20

Pandas allows you to create a data frame from a dictionary of series where you get the union of all series indices passed. As shown in Listing 6-12 with the student Salwa, no Test1 value is given. That's why NaN is set automatically.

**Listing 6-12.** Creating a Data Frame from a Series

```
In [16]: import pandas as pd
        data = {'Test1' : pd.Series([70, 55, 89],
                                   index=['Ahmed', 'Omar', 'Ali']),
               'Test2' : pd.Series([56, 82, 77, 65],
                                   index=['Ahmed', 'Omar', 'Ali', 'Salwa'])}

        df1 = pd.DataFrame(data)
        print (df1)
```

	Test1	Test2
Ahmed	70.0	56
Ali	89.0	77
Omar	55.0	82
Salwa	NaN	65

## Updating and Accessing a Data Frame's Column Selection

You can select a specific column using the column labels. For example, `df1['Test2']` is used to select only the column labeled `Test2` in the data frame, while `df1[:]` is used to display all the columns and all the rows, as shown in Listing 6-13.

### **Listing 6-13.** Data Frame Column Selection

```
In [51]: import pandas as pd
        data = {'Test1' : pd.Series([70, 55, 89],
                                   index=['Ahmed', 'Omar', 'Ali']),
               'Test2' : pd.Series([56, 82, 77, 65],
                                   index=['Ahmed', 'Omar', 'Ali', 'Salwa'])}

        df1 = pd.DataFrame(data)
        print (df1['Test2']) # Column selection
        print("\n")
        print (df1[:]) # Column selection
```

Ahmed	56
Ali	77
Omar	82
Salwa	65

Name: Test2, dtype: int64

	Test1	Test2
Ahmed	70.0	56
Ali	89.0	77
Omar	55.0	82
Salwa	NaN	65



You can select columns by using the column labels or the column index. `df1.iloc[:, [1,0]]` is used to display all rows for columns 1 and 0 starting with column 1, which refers to the column named Test2. In addition, `df1[0:4:1]` is used to display all the rows starting from row 0 up to row 3 incremented by 1, which gives all rows from 0 up to 3. See Listing 6-14.

**Listing 6-14.** Data Frame Column and Row Selection

```
In [46]: df1.iloc[:, [1,0 ]]
```

```
Out[46]:
```

	Test2	Test1
Ahmed	56	70.0
Ali	77	89.0
Omar	82	55.0
Salwa	65	NaN

```
In [39]: df1[0:4:1]
```

```
Out[39]:
```

	Test1	Test2
Ahmed	70.0	56
Ali	89.0	77
Omar	55.0	82
Salwa	NaN	65

## Column Addition

You can simply add a new column and add its values directly using a series. In addition, you can create a new column by processing the other columns, as shown in Listing 6-15.

**Listing 6-15.** Adding a New Column to a Data Frame

```
In [66]: # add a new Column
import pandas as pd
data = {'Test1' : pd.Series([70, 55, 89],
                           index=['Ahmed', 'Omar', 'Ali']),
        'Test2' : pd.Series([56, 82, 77, 65],
                           index=['Ahmed', 'Omar', 'Ali', 'Salwa'])}
df1 = pd.DataFrame(data)
print (df1)
df1['Project'] = pd.Series([90,83,67, 87],
                           index=['Ali','Omar','Salwa', 'Ahmed'])
print ("\n")
df1['Average'] = round((df1['Test1']+df1['Test2']+
df1['Project'])/3, 2)
print (df1)
```

	Test1	Test2
Ahmed	70.0	56
Ali	89.0	77
Omar	55.0	82
Salwa	NaN	65

	Test1	Test2	Project	Average
Ahmed	70.0	56	87	71.00
Ali	89.0	77	90	85.33
Omar	55.0	82	83	73.33
Salwa	NaN	65	67	NaN

## Column Deletion

You can delete any column using the `del` method. For example, `del df2['Test2']` deletes the Test2 column from the data set. In addition, you can use the `pop` method to delete a column. For example,

`df2.pop('Project')` is used to delete the column `Project`. However, you should be careful when you use the `del` or `pop` method since a reference might exist. In this case, it deletes not only from the executed data frame but also from the referenced data frame. Listing 6-16 creates the data frame `df1` and copies `df1` to `df2`.

**Listing 6-16.** Creating and Copying a Data Frame

```
In [70]: import pandas as pd
        data = {'Test1' : pd.Series([70, 55, 89],
                                   index=['Ahmed', 'Omar', 'Ali']),
               'Test2' : pd.Series([56, 82, 77, 65],
                                   index=['Ahmed', 'Omar', 'Ali', 'Salwa'])}
        print (df1)
        df2 = df1
        print ("\n")
        print (df2)
```

	Test1	Test2	Project	Average
Ahmed	70.0	56	87	71.00
Ali	89.0	77	90	85.33
Omar	55.0	82	83	73.33
Salwa	NaN	65	67	NaN

	Test1	Test2	Project	Average
Ahmed	70.0	56	87	71.00
Ali	89.0	77	90	85.33
Omar	55.0	82	83	73.33
Salwa	NaN	65	6	7 NaN

In the previous Python script, you saw how to create `df2` and assign it `df1`. In Listing 6-17, you are deleting the `Test2` and `Project` variables using the `del` and `pop` methods sequentially. As shown, both variables are deleted from both data frames `df1` and `df2` because of the reference existing between these two data frames as a result of using the assign (`=`) operator.

**Listing 6-17.** Deleting Columns from a Data Frame

```
In [71]: # Delete a column in data frame using del function
print ("Deleting the first column using DEL function:")
del df2['Test2']
print (df2)
# Delete a column in data frame using pop function
print ("\nDeleting another column using POP function:")
df2.pop('Project')
print (df2)
```

Deleting the first column using DEL function:

	Test1	Project	Average
Ahmed	70.0	87	71.00
Ali	89.0	90	85.33
Omar	55.0	83	73.33
Salwa	NaN	67	NaN

Deleting another column using POP function:

	Test1	Average
Ahmed	70.0	71.00
Ali	89.0	85.33
Omar	55.0	73.33
Salwa	NaN	NaN

```
In [72]: print (df1)
```

	Test1	Average
Ahmed	70.0	71.00
Ali	89.0	85.33
Omar	55.0	73.33
Salwa	NaN	NaN

```
In [73]: print (df2)
```

	Test1	Average
Ahmed	70.0	71.00
Ali	89.0	85.33
Omar	55.0	73.33
Salwa	NaN	NaN

To solve this problem, you can use the `df. copy()` method instead of the assign operator (`=`). Listing 6-18 shows that you deleted the variables `Test2` and `Project` using the `del()` and `pop()` methods sequentially, but only `df2` has been affected, while `df1` remains unchanged.

**Listing 6-18.** Using the Copy Method to Delete Columns from a Data Frame

```
In [83]: # add a new Column
```

```
import pandas as pd
data = {'Test1' : pd.Series([70, 55, 89],
                           index=['Ahmed', 'Omar', 'Ali']),
        'Test2' : pd.Series([56, 82, 77, 65],
                           index=['Ahmed', 'Omar', 'Ali', 'Salwa'])}
df1 = pd.DataFrame(data)
df1['Project'] = pd.Series([90,83,67, 87],
                          index=['Ali','Omar','Salwa', 'Ahmed'])
print ("\n")
df1['Average'] = round((df1['Test1']+df1['Test2']+df1
['Project'])/3, 2)
print (df1)
print ("\n")
df2= df1.copy() # copy df1 into df2 using copy() method
print (df2)
#delete columns using del and pop methods
del df2['Test2']
```

```
df2.pop('Project')
print ("\n")
print (df1)
print ("\n")
print (df2)
```

	Test1	Test2	Project	Average
Ahmed	70.0	56	87	71.00
Ali	89.0	77	90	85.33
Omar	55.0	82	83	73.33
Salwa	NaN	65	67	NaN

	Test1	Test2	Project	Average
Ahmed	70.0	56	87	71.00
Ali	89.0	77	90	85.33
Omar	55.0	82	83	73.33
Salwa	NaN	65	67	NaN

	Test1	Test2	Project	Average
Ahmed	70.0	56	87	71.00
Ali	89.0	77	90	85.33
Omar	55.0	82	83	73.33
Salwa	NaN	65	67	NaN

	Test1	Average
Ahmed	70.0	71.00
Ali	89.0	85.33
Omar	55.0	73.33
Salwa	NaN	NaN

## Row Selection

In Listing 6-19, you are selecting the second row for student Omar. Also, you use the slicing methods to retrieve rows 2 and 3.

**Listing 6-19.** Retrieving Specific Rows

```
In [106]: # add a new Column
import pandas as pd
data = {'Test1' : pd.Series([70, 55, 89],
                           index=['Ahmed', 'Omar', 'Ali']),
        'Test2' : pd.Series([56, 82, 77, 65],
                           index=['Ahmed', 'Omar', 'Ali', 'Salwa'])}
df1 = pd.DataFrame(data)
df1['Project'] = pd.Series([90,83,67, 87],index=
['Ali','Omar','Salwa', 'Ahmed'])
print ("\n")
df1['Average'] = round((df1['Test1']+df1['Test2']+df1
['Project'])/3, 2)
print (df1)
print ("\nselect iloc function to retrieve row number 2")
print (df1.iloc[2])
print ("\nslice rows")
print (df1[2:4] )
```

	Test1	Test2	Project	Average
Ahmed	70.0	56	87	71.00
Ali	89.0	77	90	85.33
Omar	55.0	82	83	73.33
Salwa	NaN	65	67	NaN

```
select  iloc function to retrieve  row number 2
Test1      55.00
Test2      82.00
Project     83.00
Average     73.33
Name: Omar, dtype: float64
```

```
slice rows
```

	Test1	Test2	Project	Average
Omar	55.0	82	83	73.33
Salwa	NaN	65	67	NaN

## Row Addition

Listing 6-20 demonstrates how to add rows to an existing data frame.

### **Listing 6-20.** Adding New Rows to the Data Frame

```
In [134 ]: import pandas as pd
           data = {'Test1' : pd.Series([70, 55, 89],
                                       index=['Ahmed', 'Omar', 'Ali']),
                   'Test2' : pd.Series([56, 82, 77, 65],
                                       index=['Ahmed', 'Omar', 'Ali', 'Salwa']),
                   'Project' : pd.Series([87, 83, 90, 67],
                                       index=['Ahmed', 'Omar', 'Ali', 'Salwa']),
                   'Average' : pd.Series([71, 73.33, 85.33, 66],
                                       index=['Ahmed', 'Omar', 'Ali', 'Salw
           data = pd.DataFrame(data)
           print (data)
           print("\n")
           df2 = pd.DataFrame([[80, 70, 90, 80]], columns
                               = ['Test1', 'Test2', 'Project', 'Average'],
                               index=['Khalid'])
           data.append(df2)
           print (data)
```

	Average	Project	Test1	Test2
Ahmed	71.00	87	70.0	56
Ali	85.33	90	89.0	77
Omar	73.33	83	55.0	82
Salwa	66.00	67	NaN	65

	Average	Project	Test1	Test2
Ahmed	71.00	87	70.0	56
Ali	85.33	90	89.0	77
Omar	73.33	83	55.0	82
Salwa	66.00	67	NaN	65
Khalid	80.00	90	80.0	70



## Row Deletion

Pandas provides the `df.drop()` method to delete rows using the label index, as shown in Listing 6-21.

### **Listing 6-21.** Deleting Rows from a Data Frame

```
In [138]: print (data)
          print ('\n')
          data = data.drop('Omar')
          print (data)
```

	Average	Project	Test1	Test2
Ahmed	71.00	87	70.0	56
Ali	85.33	90	89.0	77
Omar	73.33	83	55.0	82
Salwa	66.00	67	NaN	65
Khalid	80.00	90	80.0	70

	Average	Project	Test1	Test2
Ahmed	71.00	87	70.0	56
Ali	85.33	90	89.0	77
Salwa	66.00	67	NaN	65
Khalid	80.00	90	80.0	70

## Exploring and Analyzing a Data Frame

Pandas provides various methods for analyzing data in a data frame. The `.describe()` method is used to generate descriptive statistics that summarize the central tendency, dispersion, and shape of a data set's distribution, excluding NaN values.

```
DataFrame.describe(percentiles=None, include=None, exclude=None)
[source]
```

`DataFrame.describe()` analyzes both numeric and object series, as well as data frame column sets of mixed data types. The output will vary depending on what is provided. Listing 6-22 analyzes the Age, Salary,

Height, and Weight attributes in a data frame. It also shows the mean, max, min, standard deviation, and quantiles of all attributes. However, Salwa's Age is missing; you get the full description of Age attributes excluding Salwa's data.

**Listing 6-22.** Creating a Data Frame with Five Attributes

```
In [61]: print (df1)
data = {'Age' : pd.Series([30, 25, 44, ],
index=['Ahmed', 'Omar', 'Ali']),
'Salary' : pd.Series([25000, 17000, 30000, 12000],
index=['Ahmed', 'Omar', 'Ali',
'Height' : pd.Series([160, 154, 175, 165],
index=['Ahmed', 'Omar', 'Ali', 'Salwa'
'Weight' : pd.Series([85, 70, 92, 65], index=['Ahmed', 'Omar',
'Ali', 'Salwa']),
'Gender' : pd.Series(['Male', 'Male', 'Male', 'Female'],
index=['Ahmed', 'Omar',

data = pd.DataFrame(data)
print (data)
print("\n")
df2 = pd.DataFrame([[42, 31000, 170, 80, 'Female']], columns
=['Age', 'Salary', 'Height'
, index=['Mona']])

data = data.append(df2)
print (data)
```

	Age	Gender	Height	Salary	Weight
Ahmed	30.0	Male	160	25000	85
Ali	44.0	Male	175	30000	92
Omar	25.0	Male	154	17000	70
Salwa	NaN	Female	165	12000	65

	Age	Gender	Height	Salary	Weight
Ahmed	30.0	Male	160	25000	85
Ali	44.0	Male	175	30000	92
Omar	25.0	Male	154	17000	70
Salwa	NaN	Female	165	12000	65
Mona	42.0	Female	170	31000	80

Applying the `data.describe()` method, you get the full description of all attributes except the Gender attribute because of its string data type. You can enforce implementation of all attributes by using the `include='all'` method attribute. Also, you can apply the analysis to a specific pattern, for example, to the Salary pattern only, which finds the mean, min, max, std, and quantiles of all employees' salaries. See Listing 6-23.

### **Listing 6-23.** Analyzing a Data Frame

In [63]: `data.describe()`

Out[63]:

	Age	Height	Salary	Weight
count	4.000000	5.000000	5.000000	5.000000
mean	35.250000	144.800000	23000.000000	78.400000
std	9.215024	42.517055	8276.472679	10.968136
min	25.000000	70.000000	12000.000000	65.000000
25%	28.750000	154.000000	17000.000000	70.000000
50%	36.000000	160.000000	25000.000000	80.000000
75%	42.500000	165.000000	30000.000000	85.000000
max	44.000000	175.000000	31000.000000	92.000000

```
In [64]: data.describe(include='all')
```

```
Out[64]:
```

	Age	Gender	Height	Salary	Weight
count	4.000000	5	5.000000	5.000000	5.000000
unique	NaN	2	NaN	NaN	NaN
top	NaN	Male	NaN	NaN	NaN
freq	NaN	3	NaN	NaN	NaN
mean	35.250000	NaN	144.800000	23000.000000	78.400000
std	9.215024	NaN	42.517055	8276.472679	10.968136
min	25.000000	NaN	70.000000	12000.000000	65.000000
25%	28.750000	NaN	154.000000	17000.000000	70.000000
50%	36.000000	NaN	160.000000	25000.000000	80.000000
75%	42.500000	NaN	165.000000	30000.000000	85.000000
max	44.000000	NaN	175.000000	31000.000000	92.000000

```
In [66]: data.Salary.describe()
```

```
Out[66]: count      5.000000
mean      23000.000000
std       8276.472679
min       12000.000000
25%       17000.000000
50%       25000.000000
75%       30000.000000
max       31000.000000
Name: Salary, dtype: float64
```

Listing 6-24 includes only the numeric columns in a data frame's description.

**Listing 6-24.** Analyzing Only Numerical Patterns

In [67]: `data.describe(include=[np.number])`

Out[67]:

	Age	Height	Salary	Weight
count	4.000000	5.000000	5.000000	5.000000
mean	35.250000	144.800000	23000.000000	78.400000
std	9.215024	42.517055	8276.472679	10.968136
min	25.000000	70.000000	12000.000000	65.000000
25%	28.750000	154.000000	17000.000000	70.000000
50%	36.000000	160.000000	25000.000000	80.000000
75%	42.500000	165.000000	30000.000000	85.000000
max	44.000000	175.000000	31000.000000	92.000000

Listing 6-25 includes only string columns in a data frame's description.

**Listing 6-25.** Analyzing String Patterns Only (Gender)

In [68]: `data.describe(include=[np.object])`

Out[68]:

	Gender
count	5
unique	2
top	Male
freq	3

In [70]: `data.describe(exclude=[np.number])`

Out[70]:

	Gender
count	5
unique	2
top	Male
freq	3

You can measure overweight employee by calculating the optimal weight and comparing this with their recorded weight, as shown in Listing 6-26.

**Listing 6-26.** Checking the Weight Optimality

In [71]: data

Out[71]:

	Age	Gender	Height	Salary	Weight
Ahmed	30.0	Male	160	25000	85
Ali	44.0	Male	175	30000	92
Omar	25.0	Male	154	17000	70
Salwa	NaN	Female	165	12000	65
Mona	42.0	Female	70	31000	80

In [75]: OptimalWeight = data['Height'] - 100  
OptimalWeight

Out[75]: Ahmed      60  
Ali            75  
Omar          54  
Salwa        65  
Mona         70  
Name: Height, dtype: int64

In [93]: unOptimalCases = data['Weight'] <= OptimalWeight  
unOptimalCases

Out[93]: Ahmed      False  
Ali            False  
Omar          False  
Salwa        True  
Mona         False  
dtype: bool

## Panel Data Structures

As mentioned earlier, a *panel* is a three-dimensional data structure like a three-dimensional array.

### Creating a Panel

Pandas creates a panel using the constructor `pandas.Panel(data, items, major_axis, minor_axis, dtype, copy)`. The panel can be created from a dictionary of data frames and narrays. The data can take various forms, such as ndarray, series, map, lists, dictionaries, constants, and also another data frames.

The following Python script creates an empty panel:

```
#creating an empty panel
import pandas as pd
p = pd.Panel ()
```

Listing 6-27 creates a panel with three dimensions.

#### **Listing 6-27.** Creating a Panel with Three Dimensions

```
In [143]: # creating an empty panel
import pandas as pd
import numpy as np

data = np.random.rand(2,4,5)
Panelfdf = pd.Panel(data)
print (Panelfdf)

<class 'pandas.core.panel.Panel'>
Dimensions: 2 (items) x 4 (major_axis) x 5 (minor_axis)
Items axis: 0 to 1
Major_axis axis: 0 to 3
Minor_axis axis: 0 to 4
```

## Accessing Data from a Panel with a Position

Listing 6-28 creates a panel and fills it with random data, where the first item in the panel is a 4x3 array and the second item is a 4x2 array of random values. For the Item2 column, two values are NaN since its dimension is 4x2. You can also access data from a panel using item labels, as shown in Listing 6-28.

### **Listing 6-28.** Selecting and Displaying Panel Items

In [147]: # creating an empty panel

```
import pandas as pd
import numpy as np
data = {'Item1' : pd.DataFrame(np.random.randn(4, 3)),
        'Item2' : pd.DataFrame(np.random.randn(4, 2))}
Paneldf = pd.Panel(data)
print (Paneldf['Item1'])
print ("\n")
print (Paneldf['Item2'])
```

	0	1	2
0	-1.069595	0.835842	0.950269
1	1.063784	0.520086	1.342309
2	-2.236069	0.229717	0.752612
3	1.014550	0.903234	2.011993

	0	1	2
0	-1.126333	1.528085	NaN
1	-1.255712	0.076873	NaN
2	1.593704	-0.648342	NaN
3	0.287446	1.591275	NaN

Python displays the panel items in a data frame with two dimensions, as shown previously. Data can be accessed using the method `panel.`

`major_axis(index)` and also using the method `panel.minor_axis(index)`. See Listing 6-29.



**Listing 6-29.** Selecting and Displaying a Panel with Major and Minor Dimensions

```
In [149]: print (Panelfdf.major_xs(1))
```

```

      Item1      Item2
0  1.063784 -1.255712
1  0.520086  0.076873
2  1.342309      NaN
```

```
In [150]: print (Panelfdf.minor_xs(1))
```

---

```

      Item1      Item2
0  0.835842  1.528085
1  0.520086  0.076873
2  0.229717 -0.648342
3  0.903234  1.591275
```

## Exploring and Analyzing a Panel

Once you have a panel, you can make statistical analysis on the maintained data. In Listing 6-30, you can see two groups of employees, each of which has five attributes maintained in a panel called P. You implement the `.describe()` method for Group1, as well as for the Salary attribute in this group.

**Listing 6-30.** Panel Analysis

```
In [104]: import pandas as pd
data1 = {'Age' : pd.Series([30, 25, 44, ], index=['Ahmed',
'Omar', 'Ali']),
'Salary' : pd.Series([25000, 17000, 30000, 12000],
index=['Ahmed', 'Omar', 'Ali', 'Salwa']),
'Height' : pd.Series([160, 154, 175, 165], index=['Ahmed',
'Omar', 'Ali', 'Salwa'])}
```

## CHAPTER 6 DATA EXPLORING AND ANALYSIS

```
'Weight' : pd.Series([85, 70, 92, 65], index=['Ahmed', 'Omar',
'Ali', 'Salwa']),
'Gender' : pd.Series(['Male', 'Male', 'Male', 'Female'],
index=['Ahmed', 'Omar', 'Ali', 'Salwa'])}
```

```
data2 = {'Age' : pd.Series([24, 19, 33,25 ], index=['Ziad',
'Majid', 'Ayman', 'Ahlam']),
'Salary' : pd.Series([17000, 7000, 22000, 21000],
index=['Ziad', 'Majid', 'Ayman', 'Ahlam']),
'Height' : pd.Series([170, 175, 162, 177], index=['Ziad',
'Majid', 'Ayman', 'Ahlam']),
'Weight' : pd.Series([77, 84, 74, 90], index=['Ziad', 'Majid',
'Ayman', 'Ahlam']),
'Gender' : pd.Series(['Male', 'Male', 'Male', 'Female'],
index=['Ziad', 'Majid', 'Ayman', 'Ahlam'])}
```

```
data = {'Group1': data1, 'Group2': data2}
p = pd.Panel(data)
```

In [106]: p['Group1'].describe()

Out[106]:

	Age	Gender	Height	Salary	Weight
count	3.0	4	4.0	4.0	4.0
unique	3.0	2	4.0	4.0	4.0
top	30.0	Male	175.0	30000.0	70.0
freq	1.0	3	1.0	1.0	1.0

In [107]: p['Group1']['Salary'].describe()

```
Out[107]: count          4.0
unique          4.0
top          30000.0
freq           1.0
Name: Salary, dtype: float64
```

# Data Analysis

As indicated earlier, Pandas provides numerous methods for data analysis. The objective in this section is to get familiar with the data and summarize its main characteristics. Also, you can define your own methods for specific statistical analyses.

## Statistical Analysis

Most of the following statistical methods were covered earlier with practical examples of the three main data collections: series, data frames, and panels.

- `df.describe()`: Summary statistics for numerical columns
- `df.mean()`: Returns the mean of all columns
- `df.corr()`: Returns the correlation between columns in a data frame
- `df.count()`: Returns the number of non-null values in each data frame column
- `df.max()`: Returns the highest value in each column
- `df.min()`: Returns the lowest value in each column
- `df.median()`: Returns the median of each column
- `df.std()`: Returns the standard deviation of each column

Listing 6-31 creates a data frame with six columns and ten rows.

### **Listing 6-31.** Creating a Data Frame

```
In [11]: import pandas as pd
import numpy as np
```

```

Number = [1,2,3,4,5,6,7,8,9,10]
Names = ['Ali Ahmed','Mohamed Ziad','Majid Salim','Salwa
Ahmed', 'Ahlam Mohamed', 'Omar Ali', 'Amna Mohammed','Khalid
Yousif', 'Safa Humaid', 'Amjad Tayel']
City = ['Fujairah','Dubai','Sharjah','AbuDhabi','Fujairah','Dub
ai', 'Sharja ', 'AbuDhabi','Sharjah','Fujairah']
columns = ['Number', 'Name', 'City' ]
dataset= pd.DataFrame({'Number': Number , 'Name': Names,
'City': City}, columns = columns )
Gender= pd.DataFrame({'Gender': ['Male','Male','Male','Female',
'Female', 'Male', 'Female', 'Male','Female', 'Male']})
Height = pd.DataFrame(np.random.randint(120,175, size=(12, 1)))
Weight = pd.DataFrame(np.random.randint(50,110, size=(12, 1)))
dataset['Gender']= Gender
dataset['Height']= Height
dataset['Weight']= Weight
dataset.set_index('Number')

```

Out[166]:

	Name	City	Gender	Height	Weight
Number					
1	Ali Ahmed	Fujairah	Male	131	71
2	Mohamed Ziad	Dubai	Male	153	74
3	Majid Salim	Sharjah	Male	145	104
4	Salwa Ahmed	AbuDhabi	Female	173	86
5	Ahlam Mohamed	Fujairah	Female	158	82
6	Omar Ali	Dubai	Male	134	89
7	Amna Mohammed	Sharjah	Female	136	93
8	Khalid Yousif	AbuDhabi	Male	128	98
9	Safa Humaid	Sharjah	Female	162	81
10	Amjad Tayel	Fujairah	Male	160	77

The Python script and examples in Listing 6-32 show the summary of height and weight variables, the mean values of height and weight, the correlation between the numerical variables, and the count of all records in the data set. The correlation coefficient is a measure that determines the degree to which two variables' movements are associated. The most common correlation coefficient, generated by the Pearson correlation, may be used to measure the linear relationship between two variables. However, in a nonlinear relationship, this correlation coefficient may not always be a suitable measure of dependence. The range of values for the correlation coefficient is -1.0 to 1.0. In other words, the values cannot exceed 1.0 or be less than -1.0, whereby a correlation of -1.0 indicates a perfect negative correlation, and a correlation of 1.0 indicates a perfect positive correlation. The correlation coefficient is denoted as  $r$ . If its value greater than zero, it's a positive relationship; while if the value is less than zero, it's a negative relationship. A value of zero indicates that there is no relationship between the two variables.

As shown, there is a weak negative correlation (-0.301503) between the height and width of all members in the data set. Also, the initial stats show that the height has the highest deviation; in addition, the 75th quantile of the height is equal to 159.

### **Listing 6-32.** Summary and Statistics of Variables

```
In [186]: # Summary statistics for numerical columns
print ( dataset.describe())
```

	Number	Height	Weight
count	10.00000	10.00000	10.000000
mean	5.50000	148.00000	85.500000
std	3.02765	15.37675	10.617072
min	1.00000	128.00000	71.000000
25%	3.25000	134.50000	78.000000
50%	5.50000	149.00000	84.000000
75%	7.75000	159.50000	92.000000
max	10.00000	173.00000	104.000000

In [187]: `print (dataset.mean())` # Returns the mean of all columns

```
Number      5.5
Height     148.0
Weight      85.5
dtype: float64
```

In [188]: # Returns the correlation between columns in a DataFrame  
`print (dataset.corr())`

```
      Number      Height      Weight
Number  1.000000  0.124105  0.174557
Height  0.124105  1.000000 -0.301503
Weight  0.174557 -0.301503  1.000000
```

In [189]: # Returns the number of non-null values in each DataFrame column  
`print (dataset.count())`

```
Number      10
Name        10
City        10
Gender      10
Height      10
Weight      10
dtype: int64
```

```
In [190]: # Returns the highest value in each column
print (dataset.max())
```

---

```
Number      10
Name      Salwa Ahmed
City      Sharjah
Gender      Male
Height      173
Weight      104
dtype: object
```

```
In [191]: # Returns the lowest value in each column
print (dataset.min())
```

```
Number      1
Name      Ahlam Mohamed
City      AbuDhabi
Gender      Female
Height      128
Weight      71
dtype: object
```

```
In [192]: # Returns the median of each column
print (dataset.median())
```

---

```
Number      5.5
Height      149.0
Weight      84.0
dtype: float64
```

```
In [193]: # Returns the standard deviation of each column
print (dataset.std())
```

```
Number      3.027650
Height      15.376750
Weight      10.617072
dtype: float64
```

## Data Grouping

You can split data into groups to perform more specific analysis over the data set. Once you perform data grouping, you can compute summary statistics (aggregation), perform specific group operations (transformation), and discard data with some conditions (filtration). In Listing 6-33, you group data using City and find the count of genders per city. In addition, you group the data set by city and display the results, where for example rows 1 and 5 are people from Dubai. You can use multiple grouping attributes. You can group the data set using City and Gender. The retrieved data shows that, for instance, Fujairah has females (row 4) and males (rows 0 and 9).

### *Listing 6-33.* Data Grouping

```
In [3]: dataset.groupby('City')['Gender'].count()
```

The following output shows that we have 2 students from Abu Dhabi, 2 from Dubai, 3 from Fujairah and 3 from Sharjah grouped by gender.

```
Out[3]: City
AbuDhabi    2
Dubai        2
Fujairah     3
Sharjah      3
Name: Gender, dtype: int64
```

```
In [4]: print (dataset.groupby('City').groups)
```

```
{'AbuDhabi': Int64Index([3, 7], dtype='int64'), 'Dubai': Int64Index([1, 5], dtype='int64'), 'Fujairah': Int64Index([0, 4, 9], dtype='int64'), 'Sharjah': Int64Index([2, 6, 8], dtype='int64')}
```

```
In [5]: print (dataset.groupby(['City', 'Gender']).groups)
```

```
{('AbuDhabi', 'Female'): Int64Index([3], dtype='int64'), ('AbuDhabi', 'Male'): Int64Index([7], dtype='int64'), ('Dubai', 'Male'): Int64Index([1, 5], dtype='int64'), ('Fujairah', 'Female'): Int64Index([4], dtype='int64'), ('Fujairah', 'Male'): Int64Index([0, 9], dtype='int64'), ('Sharjah', 'Female'): Int64Index([6, 8], dtype='int64'), ('Sharjah', 'Male'): Int64Index([2], dtype='int64')}
```



## Iterating Through Groups

You can iterate through a specific group, as shown in Listing 6-34. When you iterate through the gender, it should be clear that by default the `groupby` object has the same name as the group name.

### **Listing 6-34.** Iterating Through Grouped Data

```
In [7]: grouped = dataset.groupby('Gender')
        for name,group in grouped:
            print (name)
            print (group)
            print ("\n")
```

```
Female
  Number      Name      City Gender  Height  Weight
3      4  Salwa Ahmed AbuDhabi Female   125     57
4      5  Ahlam Mohamed Fujairah Female   170     99
6      7  Amna Mohammed Sharjah Female   160     97
8      9   Safa Humaid Sharjah Female   138     70
```

```
Male
  Number      Name      City Gender  Height  Weight
0      1   Ali Ahmed Fujairah Male    130     72
1      2 Mohamed Ziad Dubai Male    129     61
2      3 Majid Salim Sharjah Male    153     51
5      6   Omar Ali Dubai Male    135     97
7      8 Khalid Yousif AbuDhabi Male    170     55
9     10 Amjad Tayel Fujairah Male    163     88
```

You can also select a specific group using the `get_group()` method, as shown in Listing 6-35 where you group data by gender and then select only females.

**Listing 6-35.** Selecting a Single Group

```
In [9]: grouped = dataset.groupby('Gender')
        print (grouped.get_group('Female'))
```

	Number	Name	City	Gender	Height	Weight
3	4	Salwa Ahmed	AbuDhabi	Female	125	57
4	5	Ahlam Mohamed	Fujairah	Female	170	99
6	7	Amna Mohammed	Sharjah	Female	160	97
8	9	Safa Humaid	Sharjah	Female	138	70

## Aggregations

Aggregation functions return a single aggregated value for each group. Once the groupby object is created, you can implement various functions on the grouped data. In Listing 6-36, you calculate the mean and size of height and weight for both males and females. In addition, you calculate the summation and standard deviations for both patterns of males and females.

**Listing 6-36.** Data Aggregation

```
In [18]: # Aggregation
        grouped = dataset.groupby('Gender')
        print (grouped['Height'].agg(np.mean))
        print ("\n")
        print (grouped['Weight'].agg(np.mean))
        print ("\n")
        print (grouped.agg(np.size))
        print ("\n")
        print (grouped['Height'].agg([np.sum, np.mean,
        np.std]))
```

```

Gender
Female    145.250000
Male      159.333333
Name: Height, dtype: float64

```

```

Gender
Female    88.750000
Male      83.666667
Name: Weight, dtype: float64

```

	Number	Name	City	Height	Weight
Gender					
Female	4	4	4	4	4
Male	6	6	6	6	6

	sum	mean	std
Gender			
Female	581	145.250000	7.274384
Male	956	159.333333	8.891944

## Transformations

Transformation on a group or a column returns an object that is indexed the same size as the one being grouped. Thus, the transform should return a result that is the same size as that of a group chunk. See Listing 6-37.

### *Listing 6-37.* Creating the Index

```

In [26]: dataset = dataset.set_index(['Number'])
         print (dataset)

```

	Name	City	Gender	Height	Weight
Number					
1	Ali Ahmed	Fujairah	Male	155	65
2	Mohamed Ziad	Dubai	Male	165	59
3	Majid Salim	Sharjah	Male	159	82
4	Salwa Ahmed	AbuDhabi	Female	138	106
5	Ahlam Mohamed	Fujairah	Female	152	100
6	Omar Ali	Dubai	Male	145	108
7	Amna Mohammed	Sharjah	Female	151	67
8	Khalid Yousif	AbuDhabi	Male	171	96
9	Safa Humaid	Sharjah	Female	140	82
10	Amjad Tayel	Fujairah	Male	161	92

In Listing 6-38, you group data by Gender, then implement the function `lambda x: (x - x.mean()) / x.std()*10`, and display results for both height and weight. The lambda operator or lambda function is a way to create a small anonymous function, i.e., a function without a name. This function is throwaway function; in other words, it is just needed where it has been created.

**Listing 6-38.** Transformation

```
In [28]: grouped = dataset.groupby('Gender')
        score = lambda x: (x - x.mean()) / x.std()*10
        print (grouped.transform(score))
```

Number	Height	Weight
1	-4.873325	-9.911893
2	6.372810	-13.097858
3	-0.374871	-0.884990
4	-9.966479	9.730865
5	9.279136	6.346216
6	-16.119460	12.920860
7	7.904449	-12.269352
8	13.120491	6.548929
9	-7.217106	-3.807730
10	1.874356	4.424952

## Filtration

Python provides direct filtering for data. In Listing 6-39, you applied filtering by city, and the return cities appear more than three times in the data set.

**Listing 6-39.** Filtration

```
In [30]: print (dataset.groupby('City').filter(lambda x: len(x)
>= 3))
```

Number	Name	City	Gender	Height	Weight
1	Ali Ahmed	Fujairah	Male	155	65
3	Majid Salim	Sharjah	Male	159	82
5	Ahlam Mohamed	Fujairah	Female	152	100
7	Amna Mohammed	Sharjah	Female	151	67
9	Safa Humaid	Sharjah	Female	140	82
10	Amjad Tayel	Fujairah	Male	161	92

## Summary

This chapter covered how to explore and analyze data in different collection structures. Here is a list of what you just studied in this chapter:

- How to implement Python techniques to explore and analyze a series of data, create a series, access data from series with the position, and apply statistical methods on a series.
- How to explore and analyze data in a data frame, create a data frame, and update and access data. This included column and row selection, addition, and deletion, as well as applying statistical methods on a data frame.
- How to apply statistical methods on a panel to explore and analyze its data.
- How to apply statistical analysis on the derived data from implementing Python data grouping, iterating through groups, aggregations, transformations, and filtration techniques.

The next chapter will cover how to visualize data using numerous plotting packages and much more.

## Exercises and Answers

- A. Create a data frame called `df` from the following tabular data dictionary that has these index labels: `['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j']`.

	Animal	Age	Priority	Visits
a	cat	2.5	yes	1
b	cat	3.0	yes	3
c	snake	0.5	no	2
d	dog	NaN	yes	3
e	dog	5.0	no	2
f	cat	2.0	no	3
g	snake	4.5	no	1
h	cat	NaN	yes	1
i	dog	7.0	no	2
j	dog	3.0	no	1

**Answer:**

You should import both the Pandas and Numpy libraries.

```
import numpy as np
import pandas as pd
```

You must create a dictionary and list of labels and then call the data frame method and assign the labels list as an index, as shown in Listing 6-40.

**Listing 6-40.** Creating a Tabular Data Frame

```
In [5]: import numpy as np
import pandas as pd
import matplotlib as mpl
```

```

data = { 'Animal': ['cat', 'cat', 'snake', 'dog', 'dog',
                    'cat', 'snake', 'cat', 'dog', 'dog'],
         'Age': [2.5, 3, 0.5, np.nan, 5, 2, 4.5, np.nan, 7, 3],
         'Visits': [1, 3, 2, 3, 2, 3, 1, 1, 2, 1],
         'Priority': ['yes', 'yes', 'no', 'yes', 'no', 'no', 'no',
                     'yes', 'no', 'no']}

labels = ['a', 'b', 'c', 'd', 'e', 'f', 'g', 'h', 'i', 'j']

#Create a DataFrame df from this dictionary data which has the
index labels.
df = pd.DataFrame( data, index = labels, columns=['Animal',
         'Age', 'Priority', 'Visits'])
print (df)

```

	Animal	Age	Priority	Visits
a	cat	2.5	yes	1
b	cat	3.0	yes	3
c	snake	0.5	no	2
d	dog	NaN	yes	3
e	dog	5.0	no	2
f	cat	2.0	no	3
g	snake	4.5	no	1
h	cat	NaN	yes	1
i	dog	7.0	no	2
j	dog	3.0	no	1

- B. Display a summary of the data frame's basic information.

You can use `df.info()` and `df.describe()` to get a full description of your data set, as shown in Listing 6-41.

**Listing 6-41.** Data Frame Summary

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Index: 10 entries, a to j
Data columns (total 4 columns):
Animal      10 non-null object
Age         8 non-null float64
Priority     10 non-null object
Visits      10 non-null int64
dtypes: float64(1), int64(1), object(2)
memory usage: 400.0+ bytes
```

In [7]: `df.describe()`

	Age	Visits
count	8.000000	10.000000
mean	3.437500	1.900000
std	2.007797	0.875595
min	0.500000	1.000000
25%	2.375000	1.000000
50%	3.000000	2.000000
75%	4.625000	2.750000
max	7.000000	3.000000

- C. Return the first three rows of the data frame `df`.

Listing 6-42 shows the use of `df.iloc[:3]` and `df.head(3)` to retrieve the first `n` rows of the data frame.



**Listing 6-42.** Selecting a Specific n Rows

```
In [12]: df.head(3)
```

```
Out[12]:
```

	Animal	Age	Priority	Visits
<b>a</b>	cat	2.5	yes	1
<b>b</b>	cat	3.0	yes	3
<b>c</b>	snake	0.5	no	2

```
In [13]: df.iloc[:3]
```

```
Out[13]:
```

	Animal	Age	Priority	Visits
<b>a</b>	cat	2.5	yes	1
<b>b</b>	cat	3.0	yes	3
<b>c</b>	snake	0.5	no	2

- D. Select just the animal and age columns from the data frame `df`.

The Python data frame `loc()` method is used to retrieve the specific pattern `df.loc[ : , ['Animal', 'Age']]`. In addition, an array form retrieval can be used too with `df[['Animal', 'Age']]`. See Listing 6-43.

**Listing 6-43.** Slicing Data Frame

```
In [16]: df.loc[:,['Animal', 'Age']]
# or
df [['Animal', 'Age']]
```

Out[16]:

	Animal	Age
a	cat	2.5
b	cat	3.0
c	snake	0.5
d	dog	NaN
e	dog	5.0
f	cat	2.0
g	snake	4.5
h	cat	NaN
i	dog	7.0
j	dog	3.0

E. Count the visit priority per animal.

```
In [8]: df.groupby('Priority')['Animal'].count()
```

F. Find the mean of the animals' ages.

```
In [10]: df.groupby('Animal')['Age'].mean()
```

G. Display a summary of the data set. See Listing 6-44.

#### **Listing 6-44.** Data Set Summary

```
In [13]: df.groupby('Animal')['Age'].describe()
```

Out[13]:

	count	mean	std	min	25%	50%	75%	max
Animal								
cat	3.0	2.5	0.500000	2.0	2.25	2.5	2.75	3.0
dog	3.0	5.0	2.000000	3.0	4.00	5.0	6.00	7.0
snake	2.0	2.5	2.828427	0.5	1.50	2.5	3.50	4.5